

# anomaly-detection-ACmotor

January 25, 2023

## 0.1 Problem Statement

To create a Machine Learning model to detect anomalies in a dataset consisting real time current readings of a 3-phase AC motor (3.2 hp)

**Approach taken to solve this problem:** The datasets are in 6 files. So combine it into a single source.

Data Analysis and Data Visualizations to understand the pattern and distribution of data

Necessary data pre-processing steps to make the data ready for modeling

As there are no target labels, this is an Unsupervised machine learning problem.

Fit the data to KMeans Clustering algorithm and split the data into 2 clusters.

The cluster value can be stored as a target column and now our problem becomes a Classification type.

Now, use the whole dataset along with target column to train and build a classification model (Random Forest Classifier)

Model is saved to a pickle file so that it can be imported and reused with any values.

The steps involved in this task: 1. Loading the tools 2. Loading the dataset 3. Exploratory Data Analysis (EDA) 4. Feature scaling 5. Model Building- KMeans Clustering algorithm 6. Model Evaluation 7. Save the model to a pickle file

## 0.2 1. Loading the tools

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
```

The data is spread across 6 different files. Combining all 6 files to a single file

```
[2]: import os
os.listdir()
```

```
[2]: ['.ipynb_checkpoints',  
      'anomaly-detection-ACmotor.ipynb',  
      'dataset',  
      'predictions']
```

```
[3]: os.chdir("dataset")  
os.listdir()
```

```
[3]: ['data0 (1).txt',  
      'data109.txt',  
      'data112.txt',  
      'data88.txt',  
      'data89.txt',  
      'data9.txt']
```

```
[4]: files = [file for file in os.listdir()]  
files
```

```
[4]: ['data0 (1).txt',  
      'data109.txt',  
      'data112.txt',  
      'data88.txt',  
      'data89.txt',  
      'data9.txt']
```

```
[5]: files_df = pd.concat(map(pd.read_csv, files), ignore_index=True)  
files_df.to_csv("files.csv")
```

### 0.3 2. Loading the dataset

```
[6]: df=pd.read_csv("files.csv").T  
df.reset_index(inplace=True)  
df.columns.values[0] = "values"  
df
```

```
[6]:
```

|       | values     |
|-------|------------|
| 0     | Unnamed: 0 |
| 1     | 0          |
| 2     | 0.1        |
| 3     | 0.2        |
| 4     | 0.3        |
| ...   | ...        |
| 13590 | 120.2065   |
| 13591 | 248.1193   |
| 13592 | 120.2066   |
| 13593 | 248.1194   |
| 13594 | 248.1195   |

[13595 rows x 1 columns]

```
[7]: df['values'].str.contains('Unnamed').sum()
```

[7]: 2

As there is no text (i.e. column name) before the first comma on the first line of the csv file, Unnamed is displayed. So we have to get rid of unnamed.

```
[8]: df["current_readings"] = df["values"].str.split(": ").str[0]
df
```

```
[8]:
```

|       | values     | current_readings |
|-------|------------|------------------|
| 0     | Unnamed: 0 | Unnamed          |
| 1     | 0          | 0                |
| 2     | 0.1        | 0.1              |
| 3     | 0.2        | 0.2              |
| 4     | 0.3        | 0.3              |
| ...   | ...        | ...              |
| 13590 | 120.2065   | 120.2065         |
| 13591 | 248.1193   | 248.1193         |
| 13592 | 120.2066   | 120.2066         |
| 13593 | 248.1194   | 248.1194         |
| 13594 | 248.1195   | 248.1195         |

[13595 rows x 2 columns]

```
[9]: txt=df[df["current_readings"] == 'Unnamed'].index
txt
```

```
[9]: Int64Index([0, 10001], dtype='int64')
```

```
[10]: df.drop(txt,inplace=True)
df.drop("values",axis=1,inplace=True)
df
```

```
[10]:
```

|       | current_readings |
|-------|------------------|
| 1     | 0                |
| 2     | 0.1              |
| 3     | 0.2              |
| 4     | 0.3              |
| 5     | 0.4              |
| ...   | ...              |
| 13590 | 120.2065         |
| 13591 | 248.1193         |
| 13592 | 120.2066         |
| 13593 | 248.1194         |

```
13594          248.1195
```

```
[13593 rows x 1 columns]
```

```
[11]: df['current_readings'].str.contains('Unnamed').sum()
```

```
[11]: 0
```

```
[12]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 13593 entries, 1 to 13594
Data columns (total 1 columns):
#   Column          Non-Null Count  Dtype
---  -
0   current_readings 13593 non-null  object
dtypes: object(1)
memory usage: 212.4+ KB
```

```
[13]: df["current_readings"] = df["current_readings"].astype(float)
```

```
[14]: df.dtypes
```

```
[14]: current_readings    float64
      dtype: object
```

## 0.4 3. Exploratory Data Analysis

```
[15]: df.describe()
```

```
[15]:      current_readings
count    13593.000000
mean         52.371722
std         80.108110
min           0.000000
25%          0.405600
50%          0.711400
75%         120.200500
max         248.999000
```

Mean = 52.37

Median = 0.71

SD = 80.10

Findings:

There is very gradual increase in values till the Q2 value (Median).

Values in the Third Quartile suddenly increases and there is a wide dispersion


### 0.4.1 Line graph

```
[16]: # Visualizing only the first 500 values to view it clearly
plt.figure(dpi=300)

df[:500].plot()
plt.xlabel("Current Readings")
plt.ylabel("count")
plt.title("Current readings")
```

```
[16]: Text(0.5, 1.0, 'Current readings')
```

<Figure size 1800x1200 with 0 Axes>




### 0.4.2 Box Plot

```
[17]: # box plot
plt.figure(dpi=300)
plt.figure(figsize=(10,8))
sns.boxplot(data=df,x='current_readings');
plt.title('Box plot', fontsize = 12)
```

```
[17]: Text(0.5, 1.0, 'Box plot')
```

<Figure size 1800x1200 with 0 Axes>



## 0.5 Histogram

```
[18]: bins=np.arange(0,400,100)
plt.hist(df["current_readings"],bins,facecolor="red",edgecolor="yellow",alpha=0.5);
plt.xticks(bins)
plt.title("Current readings distribution")
plt.xlabel("Current Readings")
plt.ylabel("count")
plt.show()
```



Findings:

It can be seen that 0 to 120 is the normal range.

Values > 200 are outliers and considered anomaly

The distribution is Right Skewed (Positive Skewness)

## 0.6 4. Feature Scaling

```
[19]: # In order to fit data to the model, it has to be reshaped.
x = np.asarray(df['current_readings'])
x = x.reshape(-1,1)
x
```

```
[19]: array([[0.000000e+00],
           [1.000000e-01],
           [2.000000e-01],
           ...,
           [1.202066e+02],
           [2.481194e+02],
           [2.481195e+02]])
```

```
[20]: # Perform Standard Scaling to restrict models from being biased towards
      ↪ features having higher / lower magnitude values.
```

```
sc = StandardScaler()
x = sc.fit_transform(x)
x
```

```
[20]: array([[ -0.6537871 ],
             [-0.65253874],
             [-0.65129038],
             ...,
             [ 0.8468228 ],
             [ 2.44363373],
             [ 2.44363498]])
```

```
[21]: x.shape
```

```
[21]: (13593, 1)
```

The data is now pre-processed and ready for modelling.

## 0.7 5. Model Building

Approach for building model:

As there are no target labels, this is an Unsupervised machine learning problem.

Fit the data to KMeans Clustering algorithm and split the data into 2 clusters.

The cluster value can be stored as a target column and now our problem becomes a Classification type.

So use the whole dataset along with target column to train and build a classification model (Random Forest Classifier)

```
[22]: ## kmeans clustering algorithm

kmeans=KMeans(n_clusters=2, init='k-means++', random_state=0)
y = kmeans.fit_predict(x)
```

```
[23]: kmeans.labels_
```

```
[23]: array([0, 0, 0, ..., 1, 1, 1])
```

```
[24]: kmeans.cluster_centers_
```

```
[24]: array([[ -0.64743859],
             [ 1.3000349 ]])
```

```
[25]: # append the predicted values to a new column

df["labels"]=kmeans.labels_
df
```



```
[25]:
```

|       | current_readings | labels |
|-------|------------------|--------|
| 1     | 0.0000           | 0      |
| 2     | 0.1000           | 0      |
| 3     | 0.2000           | 0      |
| 4     | 0.3000           | 0      |
| 5     | 0.4000           | 0      |
| ...   | ...              | ...    |
| 13590 | 120.2065         | 1      |
| 13591 | 248.1193         | 1      |
| 13592 | 120.2066         | 1      |
| 13593 | 248.1194         | 1      |
| 13594 | 248.1195         | 1      |

[13593 rows x 2 columns]

The KMeans model has grouped the data into clusters 0 and 1. It can be interpreted as : > 0 = no (There is no anomaly) ; 1= yes (This is an anomaly)

```
[26]: # making a copy of df
clustered_df = df.copy()
clustered_df
```

```
[26]:
```

|       | current_readings | labels |
|-------|------------------|--------|
| 1     | 0.0000           | 0      |
| 2     | 0.1000           | 0      |
| 3     | 0.2000           | 0      |
| 4     | 0.3000           | 0      |
| 5     | 0.4000           | 0      |
| ...   | ...              | ...    |
| 13590 | 120.2065         | 1      |
| 13591 | 248.1193         | 1      |
| 13592 | 120.2066         | 1      |
| 13593 | 248.1194         | 1      |
| 13594 | 248.1195         | 1      |

[13593 rows x 2 columns]

```
[29]: # For easy readability of dataset, change the target column to categorical
clustered_df["anomaly"] = "no"
clustered_df["anomaly"] = np.where(clustered_df["labels"] == 1, "yes",
↳clustered_df["anomaly"])
clustered_df.drop("labels",axis=1,inplace=True)
clustered_df
```

```
[29]:
```

|   | current_readings | anomaly |
|---|------------------|---------|
| 1 | 0.0000           | no      |
| 2 | 0.1000           | no      |
| 3 | 0.2000           | no      |

|       |          |     |
|-------|----------|-----|
| 4     | 0.3000   | no  |
| 5     | 0.4000   | no  |
| ...   | ...      | ... |
| 13590 | 120.2065 | yes |
| 13591 | 248.1193 | yes |
| 13592 | 120.2066 | yes |
| 13593 | 248.1194 | yes |
| 13594 | 248.1195 | yes |

[13593 rows x 2 columns]

```
[30]: clustered_df.head()
```

```
[30]:   current_readings  anomaly
1           0.0         no
2           0.1         no
3           0.2         no
4           0.3         no
5           0.4         no
```

```
[31]: clustered_df.tail()
```

```
[31]:   current_readings  anomaly
13590       120.2065      yes
13591       248.1193      yes
13592       120.2066      yes
13593       248.1194      yes
13594       248.1195      yes
```

```
[32]: # export the dataset to a csv file
```

```
clustered_df.to_csv("C:/Users/Kalpagam/udemy_bootcamp/ml_projects/
↳anomaly_detection/predictions/predicted.csv")
```

Now, readings which are outliers or anomalies are indicated as yes in the “anomaly” column.

Now, the problem has reduced to a classification problem

### Data preprocessing

```
[35]: ## split the dataset as train and test
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = \
↳train_test_split(df["current_readings"],df["labels"])
x_train.shape, x_test.shape,y_train.shape,y_test.shape
```

```
[35]: ((10194,), (3399,), (10194,), (3399,))
```

```
[44]: # reshape

x_train=np.asarray(x_train).reshape(-1,1)
y_train=np.asarray(y_train).reshape(-1,1)
x_test=np.asarray(x_test).reshape(-1,1)
y_test=np.asarray(y_test).reshape(-1,1)
x_train.shape,y_train.shape,x_test.shape,y_test.shape
```

```
[44]: ((10194, 1), (10194, 1), (3399, 1), (3399, 1))
```

```
[42]: # classification model building

from sklearn.ensemble import RandomForestClassifier

# Instantiate model
model = RandomForestClassifier(n_jobs=-1)

# Fit the model
model.fit(x_train,y_train)
```

```
C:\Users\Kalpagam\AppData\Local\Temp\ipykernel_15376\1544301963.py:9:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
    model.fit(x_train,y_train)
```

```
[42]: RandomForestClassifier(n_jobs=-1)
```

```
[45]: y_pred = model.predict(x_test)
```

```
[46]: y_pred
```

```
[46]: array([0, 1, 0, ..., 0, 0, 0])
```

## 0.8 6. Model Evaluation

```
[48]: model.score(x_test,y_test)
```

```
[48]: 1.0
```

```
[49]: from sklearn.metrics import confusion_matrix, classification_report

print(confusion_matrix(y_test,y_pred))
```

```
[[2309    0]
 [    0 1090]]
```

```
[54]: print(classification_report(y_test,y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 2309    |
| 1            | 1.00      | 1.00   | 1.00     | 1090    |
| accuracy     |           |        | 1.00     | 3399    |
| macro avg    | 1.00      | 1.00   | 1.00     | 3399    |
| weighted avg | 1.00      | 1.00   | 1.00     | 3399    |

### 0.8.1 7. Save the model to a pickle file

The model is saved to a pickle file so that it can be imported and reused

```
[57]: import pickle
file = open('C:/Users/Kalpagam/udemy_bootcamp/ml_projects/anomaly_detection/
↳trained_model.pkl','wb')
pickle.dump(model,file)
```

```
[58]: model = open('C:/Users/Kalpagam/udemy_bootcamp/ml_projects/anomaly_detection/
↳trained_model.pkl','rb')
trained_model = pickle.load(model)
```

```
[66]: sample_xtest_values = np.arange(0,150,2).reshape(-1,1)
sample_xtest_values.shape
```

```
[66]: (75, 1)
```

giving some random test values to see if it classifies it

```
[67]: result = trained_model.predict(sample_xtest_values)
print("Result:",result)
```

```
Result: [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1
1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1]
```